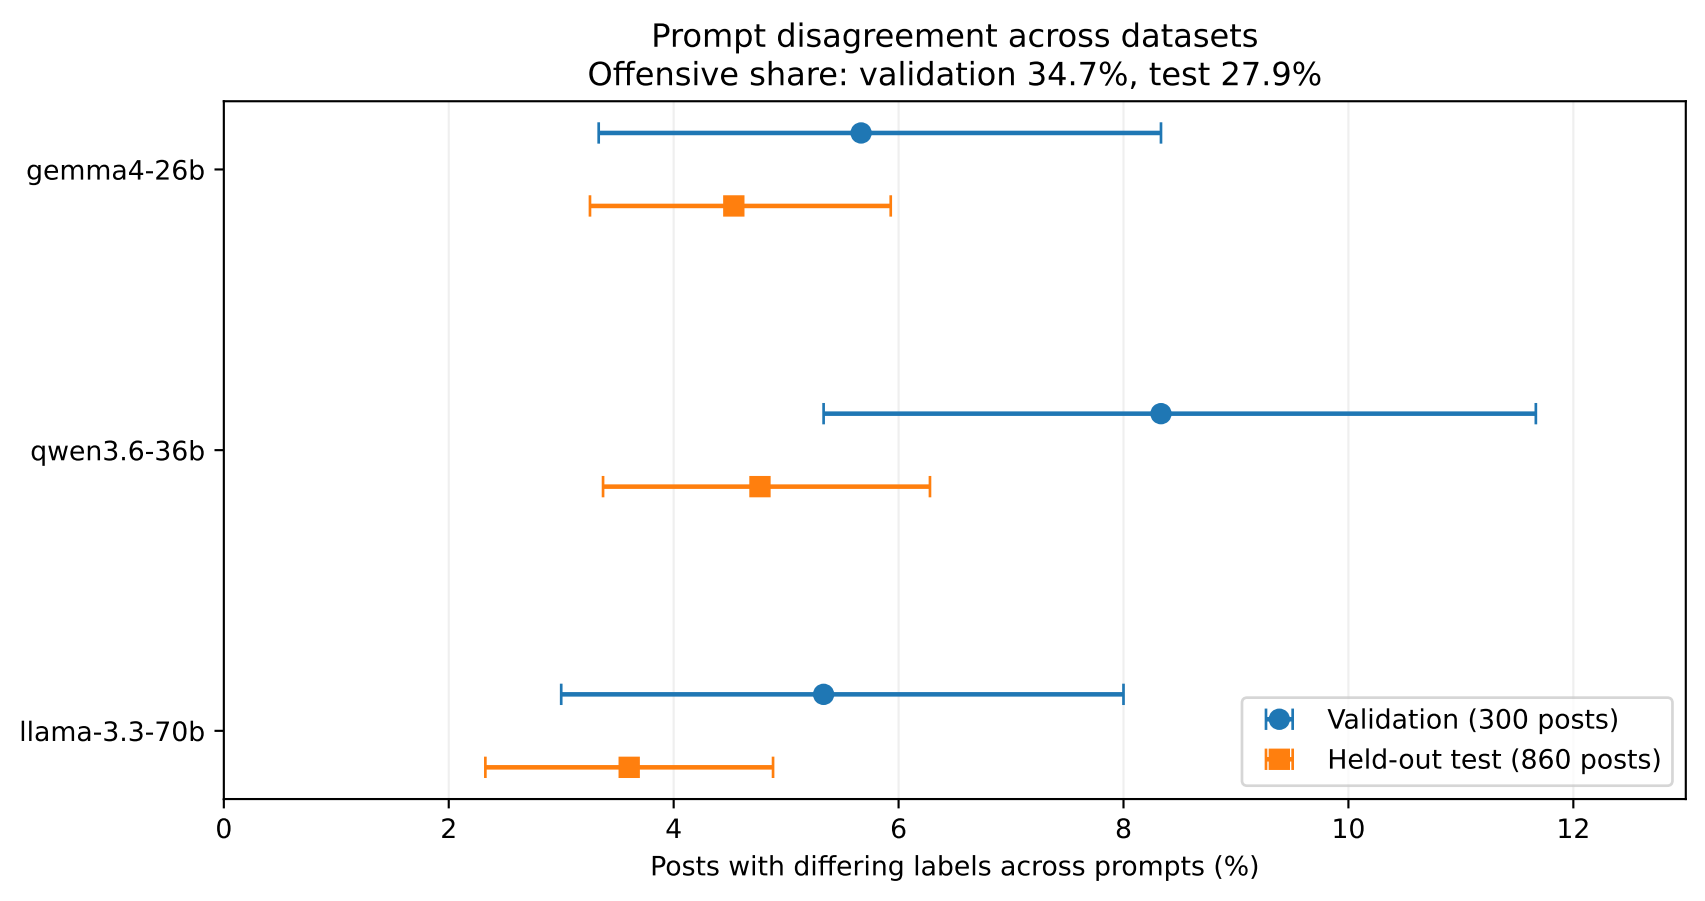
Prompt disagreement across datasets
Offensive share: validation 34.7%, test 27.9%
gemma4-26b
qwen3.6-36b
llama-3.3-70b
Validation (300 posts)
Held-out test (860 posts)
0
2
4
6
8
10
12
Posts with differing labels across prompts (%)