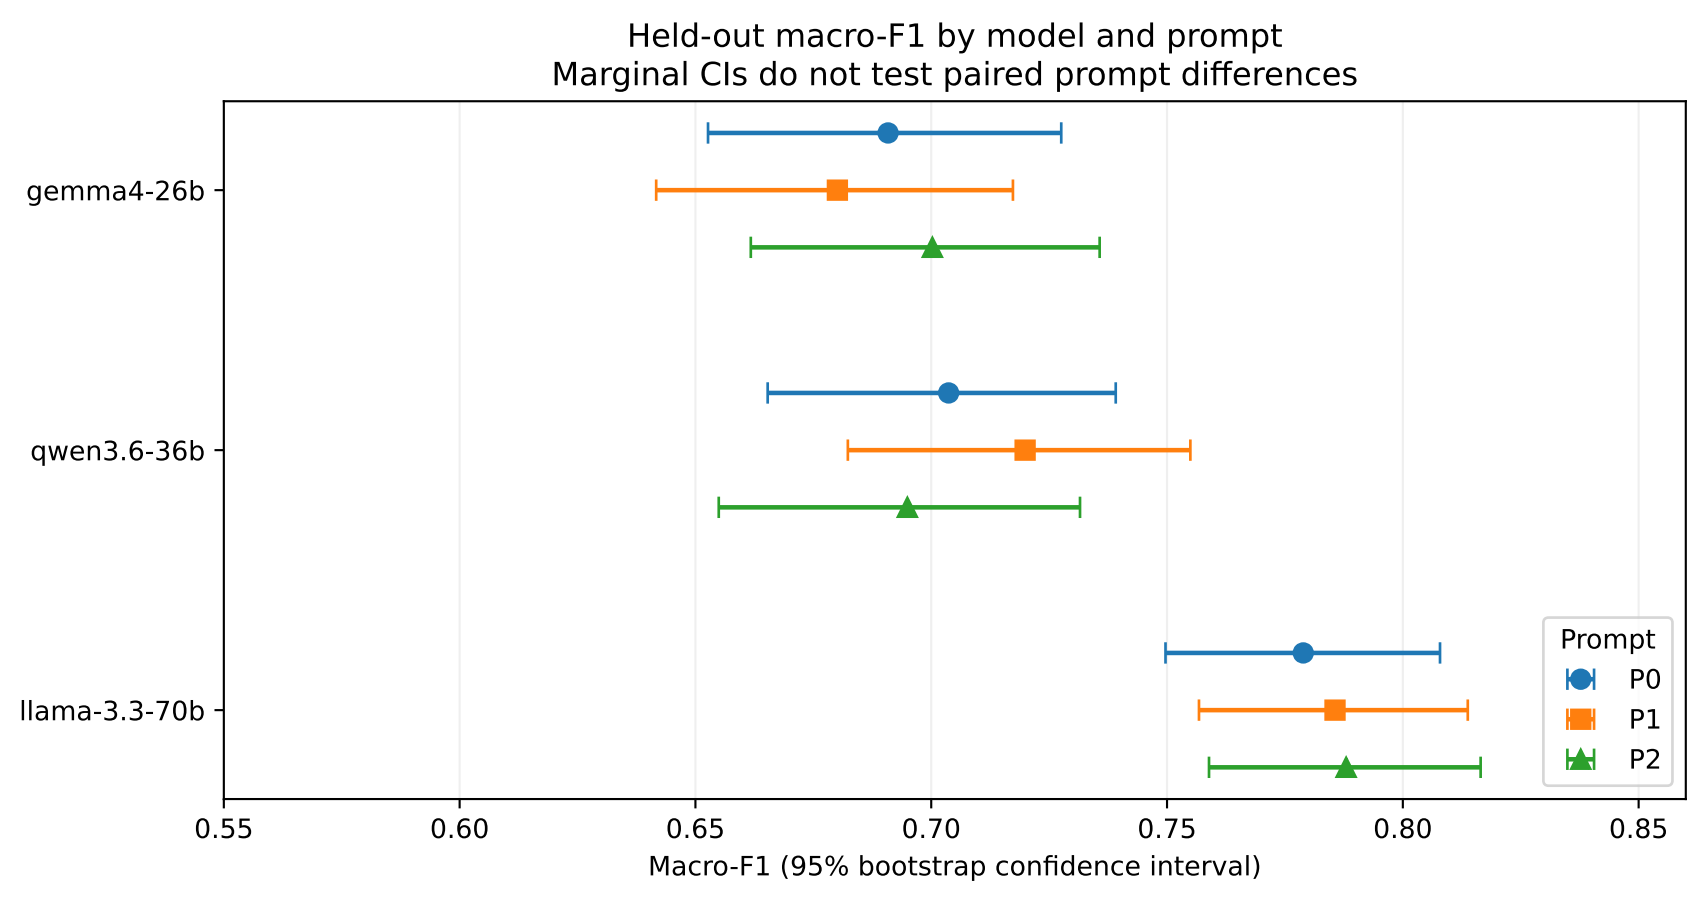
Held-out macro-F1 by model and prompt
Marginal CIs do not test paired prompt differences
gemma4-26b
qwen3.6-36b
llama-3.3-70b
Prompt
P0
P1
P2
0.55
0.60
0.65
0.70
0.75
0.80
0.85
Macro-F1 (95% bootstrap confidence interval)